

Minh · complete presentation prep

A self-contained study book: your role, 42 detailed questions, and the complete illustrated master guide.
September 26, 2026.

START HERE

Deliver the short James explanation. Take AI and label questions; Vinh adds a sentence when needed.

Reading path	What you will find
Personal page 2	Your exact spoken lines, timing, actions and next-speaker cues.
Personal pages 3–4	Questions tailored to your presentation responsibilities.
Personal pages 5–13	Shared architecture, workflow, AI, metrics, privacy, testing and production questions.
Personal pages 14–16	Advanced label identity/extraction, classifier contract and analytics ledger/math.
Master pages 1–45	The complete original guide: pitch, named team scripts, features, charts, screenshots, business scenarios and linked sources.

How to answer a judge

Give the short answer first. Expand only as needed with the implementation or evidence. Finish with the practical limit or next validation.
Each question below includes all three layers.

FIND THINGS QUICKLY

Use PDF bookmarks. P-1 to P-16 are your prep notes; M-1 to M-45 are the full master reference. Original master page numbers and links remain intact.

Source snapshot: main 015c342. Code/evidence review for these notes; application tests were not rerun for this PDF.

Minh · your speaking script

Deliver the short James explanation. Take AI and label questions; Vinh adds a sentence when needed.

2:15–2:35 / James — short sentences

“James Carter: Humira first fill confirmation is still pending.

Gemini reads the hub note.

It gives one reason: unable to reach.

If it is not sure, it leaves the reason unknown.

Gemini does not pick the fix.

A rule does.”

Action: After the opening queue, prepare James: sign ev_14, then fire ev_16–18 in order. At 2:15 Vinh shows the result; if null, explain that fallback instead.

YOUR BEST AI ANSWER

“AI reads the note. A rule picks the fix. A human taps send.”

Know these three follow-ups

- **Why not prior auth?** The earlier pharmacy reject is 75. Gemini classifies the later hub contact note, whose expected reason is unable to reach.
- **What does unknown mean?** The backend returns null for uncertain, invalid, rejected or timed-out classification. It does not invent a reason.
- **What changed?** Otezla and Humira now have verified cached labels. Humira does not gain audio or a live visit trigger from that change.

Questions: Khadim — buyer, market, partner fit. Vinh — build, data flow, polling, watch. Minh — AI and labels; Vinh supports. Keep optional features out if they risk the core proof.

Same spoken lines as master pages 4–5. Times are rehearsal targets; read actions silently.

Minh: explaining narrow, verifiable AI

Use these explanations to understand the system. Do not present a recorded test or an illustrative calculation as a new result.

Why should anyone trust the label or model output?

Short answer: They are different trust problems: labels use a source-verification contract; model output is treated as a constrained interpretation.

Do not merge the claims. The label pipeline checks cached source-derived text and identity evidence. Gemini returns only a permitted reason or null and never writes the patient narrative or chooses treatment. A model output passing the parser means it has the allowed format, not that the reason is factually correct. Demonstrate the altered-text failure for the label and safe null behavior for the classifier as separate tests.

ELI5: a verified photocopy and a reader's interpretation are not the same kind of evidence.

Is the AI robust to prompt injection or strange notes?

Short answer: The narrow output contract and deterministic router limit what an untrusted note can influence; they do not establish perfect robustness.

The source note is untrusted input. Length/blank checks, exact output parsing and allowlisted reason values reject many invalid replies. The live planner validates the command batch first and accepts the result only through server context. Injection-shaped fixtures were exercised, but a single fixture is not a broad security evaluation. A malicious or confusing note could still lead to a valid but incorrect category; real data requires adversarial and labeled quality evaluation.

Never call enum validation a proof of semantic accuracy.

Which model worked, and what do the timings prove?

Short answer: Recorded production-classifier smoke passed with gemini-3.1-flash-lite; timings are individual observations.

The handoff records an earlier gemini-3.8-flash price attempt missing the four-second deadline. A flash-lite run passed the four fixtures; later price/unreachable cases returned expected enums while null cases hit the deadline and safely returned null. That means null is intentionally ambiguous between unknown, rejection and provider failure. Do not turn four cases into an accuracy percentage or quote the fastest request as the expected production latency.

The important behavior under failure is that no scripted reason is silently restored.

Source: [scripts/labels](#) · Further module references in the master source map.

Minh: defending metrics and provenance

Use these explanations to understand the system. Do not present a recorded test or an illustrative calculation as a new result.

How do you know the SQL and displayed metrics agree?

Short answer: The project compares stored/query results with a pure summary oracle and checks replay behavior.

The pure function is an independently inspectable implementation of the metric contract: deduplicate event identities, isolate the run, establish prescription time, select a valid fill and compute the median. SQL/oracle checks cover edge cases and recorded hosted runs. Concurrent retry evidence checks stable counts. Explain which layer each test exercises; local PostgreSQL checks do not alone prove hosted Timescale setup or a rendered browser result.

The recorded hosted audit supplies the separate browser-level evidence for the live summary.

What are the awkward cases in time-to-fill?

Short answer: Duplicate events, early fills, equal timestamps, even-number medians and missing evidence all need explicit rules.

A fill before prescription does not qualify; it must not hide a later valid fill. Equal-time dispensing can qualify because prescriptions are established before fill selection. A case counts once using its earliest qualifying confirmation. The pure summary rounds the median to whole seconds; an even cohort uses the two middle values. Empty cohorts return null, not zero seconds. A confirmed fill still does not establish ingestion or treatment response.

When questioned, define the cohort and the event timestamps before interpreting the number.

What can I say about privacy and the analytics export?

Short answer: The projector minimizes fields and uses a keyed per-run case hash; that is pseudonymization and a limited data contract.

Do not describe it as anonymous or HIPAA-certified. Check what leaves each boundary: Tiger metric fields, Gemini note, ntfy alert text and optional xAI audio differ. Preserve the HMAC key for retained runs because a changed hash conflicts with immutable prior projections. The current shared demo session does not enforce a distinct pharma-only viewer. A production buyer interface needs its own authorization and disclosure controls.

Lead with exactly what is excluded from the Tiger projection, then acknowledge the wider system boundary.

Source: [lib/server/analytics/summary.ts](#) · Further module references in the master source map.

Commands, events and safe changes

Use these explanations to understand the system. Do not present a recorded test or an illustrative calculation as a new result.

What actually happens when someone presses Send?

Short answer: The screen requests an action; the server decides which events may be recorded.

The route authenticates the request and requires a run ID. The command executor reads the current snapshot, validates the complete transition, plans the permitted events and commits against the expected run/revision. The database locks the active-run pointer and writes the batch atomically. Other screens derive their view from the committed history. A client cannot submit an arbitrary new patient narrative or bypass the server by changing the button label.

ELI5: asking to add a line to the notebook is different from an accepted line.

What happens if two people act at once or press twice?

Short answer: The server checks state again and uses database identity/revision guards to avoid accepting incompatible plans.

Event identity is scoped by run and script ID. A repeat of an already-completed action can become a no-op; even that no-op checks run identity under the lock. A revision conflict causes a fresh snapshot and replan, for at most three attempts. The classifier result is cached during those retries. This is bounded conflict handling, not a guarantee that every command always succeeds or that every external service executes exactly once.

Example: two screens racing the same action must not manufacture two copies of its event.

What does reset mean, and why keep old runs?

Short answer: Reset switches the active run to a new identity; old committed history can still be read for replay.

The same fictional case IDs recur in demos, so case ID alone is insufficient. A command for an old run is rejected rather than applied to the new one. Screens discard stale responses using run/revision and local generation checks. Coordinator links and messages are run-scoped, so reset requires a fresh interactive approval. Reset is a coordinated shared-demo operation; it does not recall a notification already in flight.

ELI5: the notebook session changes, even if the names on the next page look familiar.

Synchronization, storage and access

Use these explanations to understand the system. Do not present a recorded test or an illustrative calculation as a new result.

Is this Supabase Realtime?

Short answer: No. Screens use authenticated HTTP polling, normally every 1.5 seconds.

The EventSource interface name is an internal abstraction, not browser Server-Sent Events. The adapter asks the Next server for snapshots and uses an ETag based on run and revision; unchanged state can return 304. It refuses outdated revisions and responses from an obsolete local generation. Polling makes this demo straightforward, but creates repeated requests and is not a zero-latency guarantee. Subscription design and scale would need a separate evaluation.

Say: Supabase stores state; authenticated polling synchronizes the screens.

What is shared and what is local?

Short answer: Live workflow history, coordinator relationships and approved patient messages are shared; contact marks stay on the browser.

Supabase holds the synthetic catalog, runs, events, notification outbox and coordinator/message records. Tiger stores a reduced analytics projection. Reached patient and Left message are local reminders, not calls or messages, and do not synchronize between devices. Mock mode also keeps local simulation behavior. A reliable explanation names the source of each item rather than saying that every visible control is shared live state.

Example: a message approval appears on another signed-in session; a local contact mark does not.

Does login make this a production healthcare system?

Short answer: It protects the private demo, but it is not real patient, practice or buyer authorization.

The demo uses a signed HttpOnly session cookie with a seven-day lifetime, SameSite restrictions and Secure behavior on HTTPS/production; server tools can use bearer authentication. Mutating session requests enforce same-origin checks. Browser database roles lack direct access to the protected tables/functions; the backend holds service credentials. These controls do not establish tenant isolation, patient consent, professional identity, a compliance assessment or a production audit/retention program.

Explain the protection that exists, then name the missing identity boundary.

Verified labels: what the claim means

Use these explanations to understand the system. Do not present a recorded test or an illustrative calculation as a new result.

What exactly is verified about the Otezla label?

Short answer: The application serves cached text verified against saved DailyMed source sections and drug identity evidence.

This is a source-verification pipeline, not generated medical advice. Saved artifacts carry provenance and a verification contract; the offline verifier checks source-derived text and the published fixture. The rendered Otezla sections were compared with that fixture in the recorded hosted audit. Explain literal extracted text accurately: byte-exact in this project does not mean the visual web page is identical to raw source XML, or that all UI statements are FDA-verified.

ELI5: we keep both the copied text and the receipt showing what it was copied from.

Why is a hash useful, and what does it not prove?

Short answer: A hash helps detect changes; provenance and identity checks explain what the source is.

If content changes, its digest should no longer match the expected artifact. A digest alone cannot tell whether the original document was authoritative, current or even the intended drug. That is why source identity, version information, extraction rules and the published artifact must be checked together. The recorded verifier rejects a one-character change. Avoid describing hashing as encryption, regulatory approval or proof that a document is always current.

If asked about freshness: this is a verified cached snapshot, not a runtime DailyMed fetch.

What about Humira and other medications?

Short answer: Both Otezla and Humira now have verified cached DailyMed label artifacts in merged source.

Humira was added with its own saved DailyMed/RxNorm evidence, boxed warning, volume-aware identity checks and nested KIT handling. Its branch-local endpoint/render checks and 728 tests in 54 files are recorded in Minh's sign-off; a new deployed Humira check was not run for this PDF. Further drugs still need independent verification. A prepared before-visit card is also separate from label support: it needs a scripted/offline event and has no automatic live appointment trigger in the inspected planner.

Two verified drugs do not imply universal drug coverage. Humira deployment must be checked separately from branch-local evidence.

Source: [lib/server/label.ts](#) · Further module references in the master source map.

Gemini and the rulebook

Use these explanations to understand the system. Do not present a recorded test or an illustrative calculation as a new result.

Why use Gemini if the rules choose the action?

Short answer: Gemini handles language interpretation; the deterministic router controls the administrative choice.

The classifier maps a fictional pharmacy/hub note into one permitted reason or null. The planner validates the complete command batch before sending the supported source note for classification. The live result is supplied as trusted server context, not accepted from the browser. Free-form generated text is not saved as patient narrative. Separating these jobs keeps a model response from directly choosing a card, sample or treatment.

ELI5: the reader sorts the note into a labeled box; the rulebook decides which task is allowed.

What if Gemini is wrong, slow or unavailable?

Short answer: Invalid, unknown, failed or timed-out responses become null; the system does not restore a scripted answer in live mode.

The facade rejects blank or over-140-code-point notes and applies a four-second deadline, including transports that ignore cancellation. The allowed output parser constrains format. A null reason creates no reason-triggered doctor alert, while Patient Details and access support remain available. Price refusal also needs an existing recorded quote. A valid but wrong allowed reason is still possible; real-world accuracy would require labeled evaluation beyond the few demo fixtures.

One passing smoke test proves a request path worked, not clinical accuracy or a latency guarantee.

How are card and sample actions constrained?

Short answer: The router requires a recognized reason, a matching rule and explicit eligibility flags.

A copay-card result requires the matching commercial-insurance rule and `copayCardEligible` exactly true. The bridge-sample path requires its own eligible flag and matching row. Missing, unknown or unmatched inputs fall back to `ACCESS_SUPPORT`. This is the demo's administrative routing contract. It does not verify an actual patient's program eligibility, adjudicate coverage or submit a request to a real partner.

A cost problem alone never proves someone qualifies for a manufacturer card.

Tiger: turning history into numbers

Use these explanations to understand the system. Do not present a recorded test or an illustrative calculation as a new result.

What data reaches Tiger, and why keep a separate projection?

Short answer: Only allowlisted metric events reach Tiger; the practice-side workflow remains the source of history.

Replay reads committed events for a specified run, projects supported events and writes the reduced metric batch. Persisted fields include run/script identifiers, a keyed case hash, time, metric kind and allowed reason. Names, raw notes, drug labels, insurance details and wrist text are not part of this analytics projection. This makes the reporting contract inspectable, but a pseudonymous metric store is not the same as a complete buyer-access or privacy system.

ELI5: the scoreboard receives the points it needs, not a copy of the whole notebook.

What do first fills and time to first fill mean here?

Short answer: A case counts when it has a qualifying pharmacy confirmation at or after its earliest prescription.

For each case, the summary finds the earliest prescription and earliest qualifying later-or-equal fill, then takes the median elapsed time across filled cases. Unfilled cases are excluded, so this is not everyone's waiting time. Empty results have a null median. The legacy response field recovered means a confirmed fill, not clinical recovery. Reason tallies use each case's latest non-null reason, including cases later filled; they are not necessarily the unresolved queue distribution.

Resource acknowledgment, message acknowledgment and a price change are not qualifying fills.

What do the displayed demo numbers actually establish?

Short answer: They establish that the reporting path computes the expected result from synthetic history.

The recorded deployed audit returned 8 seeded fills, a 60-second median and three reason counts of 2; reset returned 0 fills, a null median and empty reasons. Separate retained-run replay and duplicate checks passed. These numbers are not a patient study, conversion lift, time saved by staff or proof that FirstDose caused a fill. The implementation queries a hypertable directly; it does not use a continuous aggregate.

Lead with the source and cohort before interpreting any dashboard number.

Source: [lib/server/analytics-runtime.ts](#) · Further module references in the master source map.

Replay, freshness and privacy limits

Use these explanations to understand the system. Do not present a recorded test or an illustrative calculation as a new result.

How can replay avoid counting the same event twice?

Short answer: Stable event identity and immutable payload checks distinguish a repeat from a conflicting rewrite.

The analytics writer uses ledger/transaction logic around event identity. Identical retries are safe; an event with the same identity but changed canonical contents is a conflict, not a silent update. Hosted evidence includes concurrent duplicate replay with unchanged counts. This lets the summary catch up from retained history after a missed background attempt. It is not a promise that any arbitrary payload can be replayed or that a failed provider call succeeded.

ELI5: the same receipt is counted once; changing its amount under the same receipt number is an error.

How do you stop an old summary appearing after reset?

Short answer: Both server and browser require the summary to match the requested run and revision.

The server checks the active checkpoint before reading retained history, after the read, after replay and after querying. It uses a cloned history for the replay. The response carries matching run/revision headers; the client also compares its current generation and revision. A stale checkpoint returns 409, invalid query parameters 400, unavailable/configuration failure 503, and missing authorization 401. Failed replay is not disguised as a successful zero-filled summary.

A genuine empty result and an unavailable result are different states.

Are hashed analytics anonymous? Who else receives data?

Short answer: The hashes are pseudonymous; the complete data flow includes external providers.

Tiger receives a keyed case identifier rather than the raw case identity through the projector. A stable HMAC key is needed for consistent replay; changing it conflicts with existing immutable projections. The shared app session is not a separate pharma/buyer permission system. ntfy receives fictional alert text, Gemini receives the supported fictional note, and optional xAI receives the recorded command. The prototype does not control every provider's retention policy or establish real-patient compliance.

Do not say that names can never leave the practice: the notification and optional voice paths need their own explanation.

Source: [lib/server/replay-followup.ts](#) · Further module references in the master source map.

Watch, voice and patient audio

Use these explanations to understand the system. Do not present a recorded test or an illustrative calculation as a new result.

Can you guarantee exactly one watch buzz?

Short answer: No. We guard duplicate sends, but server acceptance and physical receipt are different events.

Workflow commits queue the relevant notification. A worker claims an outbox row before publishing and records accepted or unknown. It attempts at most two sends per invocation, and claimed/unknown rows are not automatically reclaimed after ambiguous failures. That favors avoiding repeated sends, with a possible missed delivery if the process fails at the wrong time. A reset cannot recall an in-flight request. Earlier Garmin receipt is recorded; the intended physical presentation run still needs its own proof.

A browser wrist mirror shows intended text; it does not establish that a watch received it.

What does Grok do, and what is still incomplete?

Short answer: It transcribes the supported spoken handoff and proposes a case for human confirmation.

The endpoint accepts bounded audio, calls xAI and parses a narrow command rather than allowing arbitrary tool use. It does not itself write a handoff. The UI converts unsupported browser recording formats when needed. A merged pre-confirm run check rejects proposals when it detects a different active run, but run capture happens after recording and only the case ID is handed onward. Complete recording-through-approval reset safety and human-device acceptance remain unverified; keep the button path available.

A successful synthetic-audio provider test does not prove microphone capture or a full confirmed handoff.

Are the patient messages generated live by AI?

Short answer: No. The eligible Otezla path uses approved templates and pre-generated English/Spanish audio.

After the scripted card-resend fix, the coordinator chooses a supported language and approves the message. The patient page retrieves that run-scoped approval, plays the committed ElevenLabs asset and can record a separate message acknowledgment. There is no runtime ElevenLabs generation, general chat, SMS integration or Humira audio. Browser decode/playback was checked; native pronunciation and audible physical-phone acceptance are separate.

Message acknowledged does not mean the audio was heard, the card was used, or the medicine was taken.

Demo tools and frontend evidence

Use these explanations to understand the system. Do not present a recorded test or an illustrative calculation as a new result.

What is simulated and what really runs?

Short answer: The screens, guarded backend and provider integrations run; patients, pharmacy events and partner fulfillment are simulated.

The seed is 13 background cases: 3 needing a fix, 2 waiting and 8 confirmed; Maria and James are separate, for 15 catalog cases. The simulator fires permitted scripted beats. The raw RxFill-shaped panel labels its projection as simulated vocabulary, not a certified wire message. Mock autoplay is not live pharmacy automation. The Relay Board is an audience visualization; its fictional price change is not measured savings.

Do not infer dispensing from a paid claim: the workflow's explicit simulated confirmation is separate.

What frontend work did the team actually finish?

Short answer: Merged repairs address layout, simulator wording, keyboard focus, search feedback, contrast, QR labels and navigation structure.

The audit found narrow-screen overflow and an incorrect all-events-fired message when prerequisites were merely blocked. Repairs distinguish loading/blocked/completed states, improve focus return and route focus, add an accessible skip destination and search announcements, and improve QR/table/landmark access. Recorded local verification includes 724 tests, live/mock builds and 57 axe scans. Those fixtures do not prove every production interaction or complete accessibility conformance.

Explain the user consequence: a small screen fits, a keyboard user can navigate, and operator status describes the actual state.

How strong is the testing evidence?

Short answer: It includes unit, SQL, browser and hosted-provider checks, each with a different scope.

The hosted audit after PR #39 exercised Maria and James, coordinator/message persistence, audio, independent-session reset and live Tiger summaries. Its 62 observations are route/browser/viewport checks, not 62 separate end-to-end scenarios. Frontend repairs had 44 checks each in Chromium and Firefox; WebKit's one skip-link failure was fixed and rechecked. Later merged code still needs targeted deployed acceptance. Tests were not rerun merely to generate this PDF.

Keep code presence, recorded tests, deployed observations and physical-device evidence as four separate claims.

Production, failure and measurement

Use these explanations to understand the system. Do not present a recorded test or an illustrative calculation as a new result.

What happens when a provider or network fails?

Short answer: The system exposes the missing capability and preserves the committed workflow where possible.

Gemini returns null safely. Analytics replay failure does not undo the command or block the independent notification attempt; a later summary request can replay retained history. Unavailable analytics is not reported as a zero success. A failed watch attempt does not prove physical delivery. Voice falls back to the button path. In the presentation, use the prepared phone or a clearly labeled replay if needed, and never silently describe mock output as a live provider result.

A successful core command and a successful downstream notification are separate facts.

What is required before real deployment at scale?

Short answer: Real data agreements, identity/authorization and operational acceptance are still substantial work.

The prototype has a single shared demo environment and fixed fictional identities. A production plan would need actual pharmacy/network integration and matching, tenant/role boundaries, appropriate data handling and retention, partner contracts, monitoring and support, load/latency evaluation and a defined recovery process. The active-run lock and 1.5-second polling are reasonable demo choices but are not evidence of population-scale capacity. TestFlight source also needs Mac signing/build/install acceptance.

Explain the next engineering work honestly; do not equate a deployed demo with production readiness.

How would a pilot prove useful and profitable?

Short answer: Measure net staff effort and meaningful workflow outcomes, then test whether a buyer will pay more than delivery costs.

Define eligible cases and a comparison/baseline first. Track staff minutes and touches, unresolved-case age, independently observed fills and time to fill, failures and manual overrides. Separate missing data from confirmed non-fill. Include integration, support, infrastructure and provider costs. The guide's 200 cases × 5 minutes = 16.7 hours and \$583 at \$35/hour is an illustrative capacity-value scenario, not observed savings or cash recovered.

A pilot should test the assumptions that drive the business model, not just generate attractive dashboard counts.

Label identity and the verification chain

Use these explanations to understand the system. Do not present a recorded test or an illustrative calculation as a new result.

How do you avoid attaching correct text to the wrong product?

Short answer: Check the product identity and source metadata as well as the copied narrative.

The identity selector treats saved RxNorm responses as untrusted and requires one branded product matching brand, ingredient, strength and form. It rejects ambiguous results, packs and substitution of a generic concept for the expected branded product. Token-boundary matching avoids treating 130 mg as 30 mg. The saved Otezla evidence uses RxCUI 1492746 for apremilast 30 mg oral tablet. SPL set/document/version/effective-time checks must agree with provenance and the published catalog.

A correct paragraph attached to the wrong strength still fails the intended verification contract.

Which parts of SPL are extracted, and how?

Short answer: Four selected section types are supported, with deterministic narrative extraction and no model paraphrasing.

The targets are boxed warning, indications, warnings/precautions and dosage/administration, identified by their section codes. An absent boxed warning can be null. The extractor preserves supported nested narrative in document order with explicit whitespace/newline/table rules, decodes entities once and rejects unsupported narrative rather than silently dropping it. Verification compares the original XML hash, method, section order/titles, exact extracted text and per-section hashes. It is not a claim that the whole SPL appears on the card.

Byte-exact refers to the project's extracted-text contract, not identical visual rendering of XML.

What prevents the build from verifying a newly rewritten answer?

Short answer: The prebuild verifier checks existing saved and published artifacts rather than rebuilding the object under test.

The source XML, provenance, RxNorm evidence, label JSON and verification receipt are checked together; published fixture/catalog values must match. The receipt links their hashes. Runtime compares fixed known drug IDs and bundled objects with the receipt; user input never becomes a filesystem path. Unknown drugs return null/HTTP 404. A known artifact failing receipt verification is marked unverified and the route returns 503. This avoids depending on a live DailyMed request during the demo, while leaving source refresh/change monitoring as future operational work.

Caching improves demo repeatability; it does not establish that the cached version will always be the latest.

Source: [lib/server/labels/identity.ts](#) · Further module references in the master source map.

Classifier contract and a precise fallback

Use these explanations to understand the system. Do not present a recorded test or an illustrative calculation as a new result.

What are the actual reason categories?

Short answer: Six business reasons are allowed; UNKNOWN becomes null.

The categories are DECLINED_AT_PRICE, COPAY_NOT_APPLIED, UNABLE_TO_REACH, PA_REQUIRED, NOT_COVERED and NOT_PICKED_UP_48H. The transport requests structured JSON with one reason key, temperature zero and no thinking budget. The parser rejects prose, markdown fences, arrays, extra fields, malformed JSON and invented categories. The model cannot supply a new fix or instruction through this output contract. These settings improve constraint, not guaranteed semantic correctness.

Example of an allowed shape: {"reason":"UNABLE_TO_REACH"}. That shape alone does not prove the note was interpreted correctly.

Where does the note come from, and what about retries?

Short answer: Only the supported existing pharmacy/hub evidence note is classified after full transition validation.

For the two live classification beats, the executor selects the associated source event note from the validated history/plan. The browser cannot send an arbitrary trusted reason override. Results are cached by source identity and note during database revision retries, so contention does not repeatedly call the model. Once a classification event is committed, retrying the command does not reinterpret it. Offline planner calls can still use scripted reasons, so always state the active mode.

The live null fallback must not quietly become a scripted confident answer.

Why collapse uncertainty and failure into the same null?

Short answer: The current contract chooses one safe administrative fallback instead of requiring downstream code to trust uncertain distinctions.

Missing configuration, rejected output, provider errors, UNKNOWN and timeout all yield null. The case can be handled through access support without a reason alert. This reduces failure modes in the demo, but limits observability: a null fixture does not prove the provider explicitly answered UNKNOWN. Production operations might need separate internal diagnostics and a human review workflow. Those would need design and privacy review rather than being assumed present.

The presentation should say unclassified, not claim that the AI confidently discovered there was no barrier.

Analytics ledger, SQL and interpretation

Use these explanations to understand the system. Do not present a recorded test or an illustrative calculation as a new result.

Why a separate event-key table beside the hypertable?

Short answer: Event identity must stay unique even if someone changes its timestamp.

The ordinary event_keys table keys run_id plus script_id and stores a payload checksum. The time-partitioned fill_events hypertable includes time in its primary key. A sorted transactional batch inserts a new ledger identity before a metric; on an existing identity it compares the immutable payload checksum. Changed content is a conflict and rolls back. Original timestamps are preserved during replay. The content checksum is SHA-256 over ordered fields; it is different from the keyed HMAC case identity.

The ledger prevents the same event from reappearing as a second row merely because its time changed.

Can you explain the median without jargon?

Short answer: Sort qualifying fill durations and take the middle value, or average the middle two.

For a purely mathematical example, durations 30, 60, 90 and 120 seconds have a median of 75 seconds. The implementation rounds the result to whole seconds and returns null when no qualifying fills exist. SQL uses PERCENTILE_CONT(0.5) and is checked against a pure TypeScript oracle. Equal-time fills can count; pre-prescription fills do not. The confirmed-only cohort excludes unresolved cases, so present unresolved counts/ages separately in a real pilot.

Do not use a falling median by itself to prove better access; the cohort and missing-data rules matter.

Why two databases, and is that faster?

Short answer: The implementation separates authoritative workflow history from a minimized analytics projection; speed superiority is not established.

Supabase owns accepted workflow events and Tiger serves derived metrics. Independent replay means analytics trouble does not undo a successful workflow action. The projection also makes exported fields explicit. A single database could compute these metrics; no benchmark demonstrates that two databases are necessary, cheaper or faster. Current reads are direct SQL, not continuous aggregates, and checkpoint checking is not one distributed ACID transaction across both stores.

Explain the responsibilities and tradeoff, not an unmeasured performance claim.

FirstDose

One shared story for Khadim, Vinh and Minh. September 26, 2026.

27%

of new prescriptions in Surescripts' analysis were never dispensed.

One missed fill. One accountable next step.

THE PRODUCT

A daily access-coordinator queue that connects a reported barrier, a reviewed next step and a later pharmacy fill signal. Proposed as an Ascend skill with a DocUpdate-style doctor experience.

Learn the pitch. Understand the system.

- **Say it.** Openings, a timed 3:15 script, stage directions and fallbacks.
- **Back it up.** Sourced statistics, value hypotheses and transparent business math.
- **Explain it.** Screens, diagrams, data flow, rules, AI boundaries and judge Q&A.;

Synthetic cases and pharmacy activity. Proposed partner integrations. Source: main 015c342. Both labels verified; five source clips added. Recorded browser evidence and physical-device acceptance are separate.

Problem source [A]: Surescripts analysis, January 2026; population details on page 8.

Find what you need fast

Full named script: pages 4–5. Individual scripts: 41–43. Feature explanations: 35–40.

03	20-second opening + 30-second summary
04–05	The complete 3:15 demo script
06	Operators, setup, two-minute cut and failure lines
07	Impiricus, DocUpdate, Ascend, Spark and FirstDose
08–09	The problem: evidence, scope and cost barriers
10–13	Savings, buyer, profitability and research inventory
14–16	Actual screens and the core case workflow
17–22	Architecture, events, routing, labels, privacy and metrics
23–25	Current status, API walkthrough and exact demo copy
26–27	Judge questions: product, technology and partner fit
28–30	Glossary, teach-back practice and one-page cheat sheet
31–34	Sources and latest-push evidence / deployment gates
35–40	Industry problem + complete feature explanations
41–43	Individual scripts: Khadim, Vinh and Minh
44–45	Latest script corrections and newly verified sources

How this was combined: Deem's supplied speaking prep + the revised research/technical guide. Conflicting lines were reconciled against source evidence; unsupported claims were rewritten.

The opening and the short answer

Keep the first sentence concrete. The complete speaker script follows on pages 4–5.

WALK-UP / ABOUT 10 SECONDS

“Sending a prescription does not finish the job. FirstDose connects a reported barrier to the coordinator who can act. Put this on.”

Hand the judge the watch. Offer it only if that device path is ready; never promise a buzz before receipt.

FIRST 15–20 SECONDS

“DocUpdate’s headline says: ‘The prescription was sent. The patient still never started it.’ Surescripts found 27 percent were never dispensed in its analysis. Sending a prescription does not finish the job.”

If you get only 30 seconds

“FirstDose is a daily queue for access coordinators. It connects a reported prescription barrier to a reviewed task and a later pharmacy fill signal. We propose it as an Ascend skill for DocUpdate. This is a fictional demonstration; the next proof is a practice pilot measuring staff effort and confirmed fills.”

ELI5

Sending an order is not the same as getting it. FirstDose is the shared checklist that notices a stuck step, assigns a helper, and waits for a separate receipt.

[A] Surescripts January 2026; [J] DocUpdate headline. New source details on page 45.

The master script · first half

Khadim opens; Vinh drives the demonstration; Minh has a short AI explanation. Timings are rehearsal targets.

0:00–0:15

KHADIM

The problem

“DocUpdate’s headline says: ‘The prescription was sent. The patient still never started it.’ Surescripts found 27 percent were never dispensed in its analysis. Sending a prescription does not finish the job.”

Show: No screens. Statistic scope: January 2026 new e-prescriptions to fill-reporting pharmacies.

0:15–0:30

KHADIM

The person

“Impiricus reports over a million opted-in healthcare providers. We focus on another person in the practice: the access coordinator. Their job is helping patients get the medicine the doctor prescribed.”

Show: Introduce the coordinator as the daily user; no claim of an untouched market.

0:30–0:50

KHADIM

What it is

“FirstDose is our proposed Ascend skill for DocUpdate: a staff workflow after a reported fill barrier. Spark engages on signals; FirstDose turns a barrier into an owned task. We would test Market Access payment per qualifying confirmed first fill. Vinh, show them.”

Show: Proposed integration and business model; no partner contract or real fulfillment.

0:50–1:00

VINH

Her morning

“This queue has three needing a fix, two waiting, and eight filled. These are fictional cases. Supabase stores the shared state; polling keeps the screens synchronized.”

Show: Seeded coordinator queue; simulation disclosure visible.

1:00–1:10

VINH

The prescription

“This is our DocUpdate-style prescribing demo. The Otezla label is verified against saved DailyMed source. New FirstDose tags mark our proposed additions.”

Show: New Rx and cached Otezla label; Humira is also verified in merged source.

1:10–1:25

VINH

It breaks

“The simulated pharmacy reports a barrier. Gemini reads the note into a reason: declined at price. Read the alert to me.”

Show: Template reason: Declined at price (\$410). Watch only if actually received.

1:25–1:35

VINH

One tap

“Send to my coordinator. One handoff tap after setup. The doctor first approves the coordinator relationship. Optional voice proposes the same handoff for confirmation.”

Show: Primary tap path; initial approval is separate from subsequent handoffs.

Fictional cases and simulated pharmacy/partner services. Read operator directions silently.

The master script · second half

Pause after the patient tap. Then show the separate pharmacy confirmation. Khadim closes.

1:35–1:45

VINH

The fix

“A rule picks the fix, not AI. Government coverage never gets a manufacturer copay card in our router. The coordinator reviews and sends through our Wallet stand-in.”

Show: Server-selected eligible action; this is simulated partner delivery.

1:45–2:15

VINH

The judge plays the patient

“Now you're Maria. Tap Use at pharmacy. That tap is acknowledgment; we still wait for the pharmacy. Now our operator sends a separate simulated fill confirmation. The board updates. The doctor gets two alerts: the barrier and the confirmed fill. Minh built the AI part.”

Show: Pause on pending after the tap. Optional approved Spanish message uses pre-generated ElevenLabs audio.

2:15–2:35

MINH

James — short sentences

“James Carter: Humira first fill confirmation is still pending.

Gemini reads the hub note.

It gives one reason: unable to reach.

If it is not sure, it leaves the reason unknown.

Gemini does not pick the fix.

A rule does.”

Show: Latest hub note, not the earlier PA rejection. Null leaves the case unclassified; no scripted reason or reason alert is substituted.

2:35–2:45

VINH

The proof

“The rule routes to access support. Tiger receives selected metrics: confirmed fills and time to first fill. This view shows aggregates without names.”

Show: Check actual Tiger source badge; before-visit card is a prepared optional view, not a live trigger.

2:45–3:05

KHADIM

Three proposed upgrades

“We build on the Impiricus ecosystem. DocUpdate gets a proposed fill signal and staff seat. Ascend gets a barrier-to-task workflow. Wallet, Concierge, QPharma and Medvantx are the resource destinations. Our next proof is a practice pilot measuring staff time and independently confirmed fills.”

Show: Count three proposed upgrades; partner calls remain stand-ins.

3:05–3:15

KHADIM

Close

“Impiricus reaches the doctor who writes the prescription. FirstDose reaches the person who gets it filled. One missed fill. One accountable next step.”

Show: Final timeline; no claim of medication ingestion or proven recovery.

Optional audio requires the approved message. Physical alert receipt must be observed in the run.

Run the table without losing the story

Khadim: opening, business and close. Vinh: desktop, phone, patient, watch and build questions. Minh: short AI scene, labels and classification; helps operate the simulator on cue.

Before judges arrive	Check
One known run	Check build/mode; coordinate reset. Seed the queue, then request and approve the interactive coordinator link again. Reset clears links and messages.
Three ready surfaces	Coordinator desktop, doctor phone, patient phone. Prepare the private demo login and a spare phone before handing over the QR.
Verify the proof beats	Patient tap remains pending; separate simulator confirmation updates the other device. Confirm the Otezla card, sound and actual watch path.
Prepare James after queue	Sign James's order (ev_14), then fire ev_16–18 in order before 2:15. If unready, skip or disclose a prepared replay. Null stays unclassified; use access support.

Short, honest recovery lines

- **Watch does not arrive.** “The alert is visible here; wrist delivery did not arrive in this run.” Show the phone alert.
- **QR or login stalls.** “Use our prepared patient phone for this step.” Keep the same case and run.
- **Network fails.** “I’m switching to a labeled replay/recording of the workflow.” Explain that it is a fallback.
- **Time is short.** Cut James first, then optional audio/voice and analytics detail. Keep the separate acknowledgment and confirmation.

Two-minute cut: opening → queue/barrier → handoff → resource → acknowledgment still pending → separate fill → close. A recorded three-browser check is not a current physical-device check.

Where FirstDose would fit

The diagram is our proposed product placement. It does not indicate a live Impiricus partnership.



Name	Plain-language role / boundary
DocUpdate	Prescriber-facing application. Our phone view is a FirstDose concept modeled on that experience, not the actual partner application.
Ascend	Impiricus describes real-time access to representatives and brand resources. Proposed home for this workflow; integration is simulated here.
Spark	Lists first-time prescriptions among multiple engagement triggers. FirstDose demonstrates coordinator work after a reported barrier; we do not claim Spark cannot address it.
Wallet / Concierge / partners	Named resource paths in the concept: card delivery, support and sample requests. FirstDose does not currently call real partner fulfillment APIs.

WHY THIS IS USEFUL TO A PARTNER

The hypothesis is a repeatable staff workflow around an actionable prescriber engagement. Validate the overlap and distribution opportunity with the partner; do not claim an untouched market.


Public product roles: [E] DocUpdate FAQ, [G] Impiricus products. Resource concepts: project spec [1].

The problem, in three numbers

These are external problem statistics. None is a measured FirstDose outcome.

27 out of 100

Visualizing the reported 27% unfilled share.
This is not a sample of FirstDose users.



■ Reported unfilled share

13 HOURS / WEEK

Average physician-and-staff time per physician spent completing prior authorizations in the AMA survey.

40% OF PHYSICIANS

Reported employing staff who work exclusively on prior authorization. The same survey reports 40 requests per physician per week.

What the figures do—and do not—establish

- **A substantial access gap.** The Surescripts figure covers new e-prescriptions sent to fill-reporting pharmacies in January 2026. It is not a drug-specific abandonment rate.
- **A costly administrative job.** AMA measures prior-authorization work, including medication and medical-service requests. FirstDose does not automate all of it.
- **A reason to test our product.** These findings justify a pilot. They do not establish how many prescriptions we recover or how much money we save.

[A] Surescripts, January 2026; [B] AMA survey of 1,000 practicing physicians, December 2025.

Cost is a real barrier—use the right year

A useful supporting fact, not a prediction of what FirstDose will change.

WHAT THE PRIMARY SOURCE REPORTS

IQVIA's October 2020 account reports **44% abandonment when prescription out-of-pocket cost exceeded \$125**, and **more than 60% when cost exceeded \$500**. These describe historical populations.

Cost barrier ≠ program eligibility

A price problem identifies work to review. It does not prove that a particular person qualifies for a card or sample.

How to use it in questions

- **Say the date.** “Historical IQVIA research also associates higher out-of-pocket costs with greater abandonment.”
- **Keep the groups separate.** The >\$125 and >\$500 thresholds overlap. They are not additive slices of the population.
- **Do not infer a treatment effect.** These figures do not say that reducing a demo quote from \$410 to \$0 will recover a given share of prescriptions.
- **Keep 27% as the opening.** The newer Surescripts first-fill statistic has a different population and definition; do not blend the two datasets.

Source choice: this master uses the directly checked October 2020 account's 44% and >60%. The supplied prep used a different 2020 reference for <5% / 45% / 60%; do not interchange the series.

[H] IQVIA, Michael Kleinrock, October 8, 2020. Historical prescription-cost discussion.


What could FirstDose actually save?

Potential value: less status chasing, fewer duplicated touches, and faster identification of the next owner.

A SIMPLE EXAMPLE

200 cases/month $\times$ 5 minutes less work per case $\div$ 60 = **16.7 hours/month**.

At an assumed loaded labor cost of \$35/hour: **about \$583/month of staff capacity value.**



Sensitivity only: 200 cases/month and \$35/hour held constant. The \$250/month practice fee is a hypothetical pricing test, not an announced FirstDose price.

- **Potential practice value.** At five minutes saved, \$583 less a hypothetical \$250 fee leaves \$333 in monthly capacity value before onboarding and other costs. Freed time is not automatically payroll savings.
- **Potential patient value.** Less waiting and better access to an eligible resource. The demo's \$410 $\rightarrow$ \$0 price change is fictional; copay relief does not prove lower total healthcare spending.
- **What must be measured.** Net minutes and touches per case, including new work the software creates. Confirm any reduction against a defined baseline.

All economic inputs and savings below are illustrative assumptions, not pilot results.

Yes, it is a software product

Product concept today. A commercially validated business requires evidence of adoption and willingness to pay.

Question	Concrete answer
Who uses it?	An access coordinator works the daily queue. A prescriber approves and receives relevant alerts. The patient receives the scoped resource.
What is the product?	A workflow service: barrier evidence, ownership, allowed next step, communication and subsequent-fill tracking in one shared case history.
Who might buy it?	A platform partner such as Impiricus, a patient-access program, or a practice. These are buyer hypotheses; no agreement is established.
Why pay repeatedly?	The work recurs with new prescriptions. Payment is justified only if the product adds value after integration, training and ongoing operating effort.

Two commercial hypotheses to test

- **Workflow license or subscription.** A platform contract or practice subscription could pay for ongoing use. This is a possible model, not a confirmed deal.
- **Fee per qualifying confirmed case.** Earlier project materials proposed Market Access payment per confirmed first fill. This needs a buyer-agreed definition, attribution and commercial review; a fill counter alone is not a billable outcome.

WHY THIS COULD ALSO REMAIN A FEATURE


Surescripts already sells first-fill monitoring. FirstDose's proposed contribution is the coordinator workflow around the signal. Differentiation, data access and distribution must hold up against existing products and a partner building the feature itself.

Project positioning: [1–3]. External fit and existing alternatives: [A], [E].

Could it be profitable? Test the math.

Potentially. Revenue has to cover real data, delivery, support, sales, integration and operating costs.

Assumption / calculation	Example per month
Billable volume × assumed fee	1,000 qualifying cases × \$50 = \$50,000 revenue
Variable delivery cost	1,000 billable cases × assumed \$20 = \$20,000 ; must include allocated nonbillable work
Contribution before fixed costs	\$50,000 – \$20,000 = \$30,000 (60%)
Assumed fixed operating budget	\$25,000 for engineering, sales/admin and operations
Illustrative operating surplus	\$30,000 – \$25,000 = \$5,000 , before taxes/financing



Break-even = fixed costs ÷ (fee – variable cost) = \$25,000 ÷ \$30. If the fee falls to \$40, or variable cost rises to \$30, break-even becomes 1,250 cases/month.

The assumptions are the hard part. No buyer has validated the \$50 fee; no operating study supports the \$20 or \$25,000 costs. Practice labor savings alone do not justify \$50/case. A partner needs separately demonstrated value. One-time setup and acquisition costs also need recovery.

The research behind the pitch

External facts establish context. Team calculations and commercial assumptions require separate validation.

Existing research	What the updated guide establishes
Abandonment + workload	27% Surescripts figure rechecked. AMA hours and dedicated-staff share rechecked; newer 2025 survey has 40 PAs/week, replacing the older 39.
Georgia prescribing data	CMS 2024 Part D: Otezla—614 prescribers, 6,661 claims; Humira(CF) Pen—717 prescribers, 15,295 claims. Aggregate rows re-fetched [C–D].
Product fit	DocUpdate FAQ says staff/practice accounts are on its roadmap [E]. Its fill-confirmation wording appears in a cancellation answer; do not claim a complete product gap from that alone.
National workforce model	AAMC reports 866,460 direct-care physicians [F]. The team’s ~282,000 full-time-equivalent workload estimate extrapolates survey hours; it is not coordinator headcount or sellable seats.
Coordinator market estimate	The prior ~69,000–116,000 range assumes 1 dedicated staff member per 3–5 physicians. That staffing ratio and paid addressable market are unvalidated.
Broader reference data	The repo also contains formulary, acquisition-price and payer-mix research. Those figures describe markets and datasets; they do not measure FirstDose savings.


THE MISSING EVIDENCE IS OUR OWN

We have public evidence of a problem and a prototype of a response. We still need net time saved, comparative fill outcomes, buyer willingness to pay, actual data/support costs and acquisition/retention evidence.

Repository: docs/research/public-data-sources.md; public-sources-briefing.md; data/reference/public-data.json.

Start the demo at the queue

The desktop is the coordinator’s workspace. Use the queue to explain the product before showing the stack.




Recorded hosted audit capture, September 26, after PR #39. It precedes the latest frontend repairs and is not a new deployment check.

Point to	Explain
Summary tiles	Three needing a fix, two waiting, eight fill confirmations in this seeded background scenario.
Reason + next step	The row connects the reported barrier to an administrative action.
Waiting on	An owner makes the next step clear. An unresolved case stays visible.
Coordinator approval	Approve the interactive relationship in Profile before handoff; this seeded view shows background cases.

Source: recorded September 26 hosted audit; synthetic seeded queue. Evidence scope on page 34.

The doctor gets a short handoff

A concept for FirstDose inside DocUpdate. The demo uses FirstDose’s own interface and synthetic records.



Notice three things

- **A reported status.** The alert begins with the pharmacy-shaped status, then the reason.
- **One clear action.** “Send to my coordinator” makes the next owner explicit.
- **A visible trail.** The prescription status progresses toward a separately confirmed fill.

ELI5

The phone is the “please help with this” button. The desktop is where the helper works the checklist.

Profile supports coordinator approval. The hosted audit verified it across sessions; this saved image shows the alert surface.

Saved hosted-browser evidence, not proof of current physical phone/watch delivery.

Sources: recorded hosted audit; [3] integration; [6] templates. Screenshot predates latest frontend repairs.

Walk one case through the system

Learn the event boundaries. They are more important than memorizing every screen.



THE LINE TO MEMORIZE

A resource acknowledgment is not a pharmacy fill. A pharmacy fill is not proof of a first dose or clinical improvement.


Message acknowledgment and savings-card acknowledgment are also distinct. Neither creates dispensing evidence.

Sources: [3] integration; [4] workflow; [6] patient confirmation templates.

The architecture in one picture

The implemented live adapter uses HTTP polling through the practice server.

Stack: Next.js 16 · React 19 · TypeScript · Tailwind 4 · Supabase / PostgreSQL



SUPPORTING INPUTS


The simulator supplies pharmacy-shaped events. Verified DailyMed/RxNorm artifacts support the label card. Voice is optional; Gemini classification and Tiger replay are merged. Recorded deployed browser workflows and live summary checks passed; physical-device acceptance remains separate.

"EventSource" here is a project interface name, not a claim of browser Server-Sent Events. Polling is implemented despite the filename `realtime.ts`.

Sources: [4] workflow + migrations; [7] `lib/realtime.ts` and API routes; [8] labels.

Commands, events and safe updates

An action asks for a change. An event records a change the server accepted.



ELI5: ONE SHARED NOTEBOOK

Each screen reads the same notebook. You ask the teacher to add a line; you do not rewrite the notebook yourself. If two people ask at once, the teacher checks which action is still allowed.

- **Atomic transition.** The database checks and writes the workflow change together. This limits races between simultaneous actions.
- **Idempotent behavior.** Repeated requests should not create the same scripted event twice. Uniqueness is scoped to a run and script identity.
- **Run identity.** Reset begins a new demo run. Clients clear the old run before displaying the new one.
- **Revision + stale-response guards.** Late responses must not overwrite newer state or restore a retired run.
- **Polling.** The live adapter asks the server for a snapshot on a roughly 1.5-second interval and de-duplicates delivered events.


WHY A JUDGE SHOULD CARE

These choices make the demo repeatable across browsers and keep a fast double tap or delayed request from silently changing the story.

Sources: [4] workflow and migrations; [7] polling adapter.

Rules decide. AI can help read.

Explain the classifier and router as separate components with different authority.



Check	Result in the current router
Unknown / invalid reason	Return ACCESS_SUPPORT. Do not invent a reason.
Government coverage	The configured rules lead to access support, not a manufacturer copay-card action.
Copay-card route	Requires a matching commercial-insurance rule AND explicit copayCardEligible === true.
Bridge-sample route	Requires a matching rule AND explicit bridgeSampleEligible === true.
Missing eligibility / no match	Return ACCESS_SUPPORT for review.

ELI5: READER AND RULEBOOK


AI can help put a messy note into a labeled box. The rulebook decides which action is allowed. A label on a box is not permission to give someone a benefit.

These are prototype routing rules, not a legal determination or a complete program-eligibility engine. A reason may be supplied directly without a model call.

Sources: [9] lib/server/router.ts, mock/reasons.json; [13] merged classifier integration.

How the label earns your trust

The label pipeline serves saved, verified material. It does not generate medication claims.



ELI5: CHECKED PHOTOCOPY

Think of a carefully checked copy of an instruction sheet. The copy keeps the source's words, and a fingerprint helps detect changes. The fingerprint alone cannot prove the original sheet was authentic.

Current verified-label scope: Otezla and Humira artifacts are included. Humira added strength/concentration, volume and nested KIT identity checks. "Byte exact" refers to the project's extracted-text verification contract, not a claim that the rendered page is identical to raw XML.

Sources: [8] lib/server/label.ts; labels/*; data/labels/drug_otezla/*; package prebuild.

Who sees what, and why

Separate the intended product boundary from the demo's identity and transport mechanisms.

Surface	Intended information
Practice	Case details, reported barriers, coordinator work, prescriber approval and fill status.
Patient	The resource and approved message for the selected demo case.
Partner / Market Access	Allowlisted aggregate measures; not a copy of the practice's full case history.

What the implementation actually gives you

- **Private demo access.** Practice workflow commands use a demo cookie and server-side database credentials. The summary endpoint also requires demo authentication. The shared demo login still does not isolate real patient/practice roles.
- **Persisted approval and assignment.** The live demo uses fixed synthetic identities. Changing a display name does not grant another case.
- **Shared messages; local contact marks.** Approved EN/ES messages persist by run. "Reached patient / Left message" contact marks remain device-local.
- **Watch delivery has its own proof.** An ntify HTTP acceptance is not proof that a person saw a wrist alert. Check the actual phone/watch path.


ELI5: DIFFERENT WINDOWS

People need different windows into the same work. A summary window should not expose the whole notebook. A separate partner authorization boundary and production compliance are not established by this demo.

Sources: [3] integration handoff; [10] coordinator/message routes + migrations; [11] ntify.

Read the graphs honestly

This chart describes the seeded fictional background queue. It is not a real-world success rate.



13 background cases = 3 + 2 + 8. The two interactive cases are separate, giving 15 cases in the fixture set. Counts reflect the seed state, not a completed interactive run.

What to measure in a pilot

Measure	Meaning
Time to first fill	Timestamp of first separate pharmacy confirmation minus prescription timestamp. Define the cohort and clock.
Unresolved-case age	How long an open access problem has waited. Do not silently drop unresolved cases from analysis.
Coordinator effort	Time or touches per task, compared with a defined baseline.
Confirmation count	Count each distinct qualifying case once in the chosen run/cohort; disclose source and freshness.

A seeded count of 8 is not evidence that FirstDose caused 8 fills. No causal improvement, clinical outcome or return-on-investment estimate is established here.

Sources: [12] seed-week handoff + fixtures; [3] integration handoff.

What is built vs. what is verified

Code presence, a reported test, a deployed check and a physical demonstration are different kinds of evidence.

Capability	Evidence / wording to use
Workflow + sync	Merged APIs, PostgreSQL transitions and 1.5-second polling. Recorded hosted Maria/James workflows and independent-session reset checks passed.
Coordinator + messages	Hosted approval/messages and Otezla EN/ES playback passed. Audio is pre-generated ElevenLabs; no runtime call. Physical/native Spanish review remains.
Two verified labels	Saved source, identity data and verification pipeline present. Humira is now verified too; the prebuild gate checks both drugs.
Pharmacy / partners	Simulated pharmacy/hub activity. Impiricus, DocUpdate, Wallet and sample-partner surfaces are concepts/stand-ins.
Watch	ntfy publisher implemented. Repository records prior Garmin receipt. Latest Apple Watch acceptance remains unverified here.
Voice / Grok	Pre-confirm run check merged. Full recording-through-approval reset safety and human-device acceptance remain unverified; keep voice optional.
Gemini	Merged and wired into live commands. Recorded hosted runs returned price/unreachable reasons; null fallback remains. Rules choose the action.
Tiger	Replay, auth, freshness and pharmacy-fill projection merged. Deployed audit: 8 seeded fills, median 60 seconds; reset gives 0/null. Synthetic data.
Hosted / TestFlight	Production recorded live. Physical two-device/watch run remains. Frontend fixes merged; deployed recheck separate. iOS signing/install needs a Mac.

Before presenting, check the exact deployed commit and mode. Some older pitch/status documents still describe capabilities that later code supersedes.

Fetches September 26: main 015c342. Recorded deployed evidence is distinct from this source snapshot; see page 34.

Trace an action all the way through

A screen asks; the server checks; the database commits; other screens refresh.

Route / component	What to explain
POST /api/rx	Starts the fictional prescription workflow.
POST /api/handoff	Validates the case and coordinator flow, then records the handoff.
POST /api/fix	Checks the requested action against server-side routing and case state.
POST /api/patient/use	Emits the resource acknowledgment event only. For the interactive case this is ev_10.
POST /api/sim/fire	Accepts an ordered list of scripted IDs. ev_11 is the separate simulated pharmacy confirmation after acknowledgment.
GET /api/events	Returns run identity, revision and events. The live adapter polls and rejects stale state.
Reset / coordinator / message	Reset begins a new run; coordinator approval and patient messages have their own persistence routes.
GET /api/access/summary	Authenticated replay-before-read and matching run/revision headers. Recorded live summary passed. Stale = 409; unavailable = 503.

ELI5: ASK, CHECK, WRITE, READ

Two people can press a button. The server and database decide what is allowed to become a line in the notebook. Every screen then reads the accepted story.

Sources [4], [7], [14]: workflow, live adapter and latest analytics modules.

Read the real copy; explain the tags

Use the current display text when narrating a case. Do not replace pending status with an invented clinical conclusion.

PENDING ALERT

“Maria Lopez: Otezla first fill pending”

AFTER THE RESOURCE TAP

“Maria Lopez acknowledged the Otezla savings card. Pharmacy fill confirmation is still pending.”

AFTER SEPARATE PHARMACY CONFIRMATION

“Maria Lopez: Otezla pharmacy fill confirmed. This does not confirm treatment start.”

UNRESOLVED BEFORE-VISIT VIEW

“James Carter: Humira first fill confirmation is still pending. Review fill status before the visit.”

What “New · FirstDose” means

“The tags highlight our proposed additions. The surrounding interface is our concept of the DocUpdate experience.”

The tags now exist in code. They do not turn our UI into the real partner app or prove exact feature parity. Before-visit copy needs a prepared scripted/offline event; no automatic live visit trigger. Otezla and Humira labels verified in merged source.

Source [6]: mock/templates.json at main 015c342. Values below fill existing fictional-case template variables.

Answers to likely judge questions

Give the direct answer first, then offer evidence or the next test.

“What makes this more than a dashboard?”

It assigns a next step, validates actions, preserves an event trail and distinguishes acknowledgment from pharmacy confirmation.

“Where is the AI?”

Gemini maps the pharmacy note to a reason or null. Deterministic rules choose the fix. Voice proposes a handoff for human confirmation.

“Are you connected to Impiricus or a pharmacy?”

Those partner connections are simulated concepts here. The workflow/backend is our implementation; a real partner feed is a next integration milestone.

“Does the patient button mean the medicine was taken?”

No. It records acknowledgment. A separate pharmacy event records a fill, which still does not prove a dose was taken.

“Who pays, and how much?”

A partner-sponsored workflow or pilot is a commercial hypothesis. We have no validated pricing, willingness-to-pay result or customer contract to claim.

“What is technically difficult?”

Keeping multiple views consistent while enforcing valid actions, separating runs, preventing stale updates and preserving source-backed evidence.

“What would you do next?”

Verify the latest deployed physical flow, validate the workflow with practice staff, and agree a real pharmacy/partner data integration.

Partner and evidence questions

Keep the user benefit clear while naming the integration boundary.

“Isn’t this already Spark?”

Spark lists multiple engagement triggers. Our proposed contribution is the coordinator’s barrier-to-task-to-follow-up workflow. We need partner feedback on overlap.

“Don’t you rebuild Wallet or DocUpdate?”

The demo reproduces enough UI to show the concept. We propose routing to existing resources; real partner APIs are stand-ins here.

“Where would a real fill signal come from?”

A licensed and validated pharmacy/network integration. Our simulator uses pharmacy-shaped status vocabulary; that alone does not establish network compatibility or certification.

“Does pharma see patient names?”

The intended partner view contains aggregates. The shared demo login and prototype boundaries do not establish production role isolation. Internal hashed metrics are pseudonymous, not automatically anonymous.

“Does it send SMS?”

The demonstrated patient path is a QR-opened web page. Messages/audio play there. The wrist path uses ntfy push to the paired phone/watch.

“Why use synthetic records?”

They make the demo repeatable without real patient records. Real operation needs appropriate access, privacy controls and partner agreements; HIPAA is not a blanket ban on all authorized use.

“Have you proven the economics?”

No. The model shows what must be true. We need measured net staff time, actual delivery costs and a budget owner willing to pay.

Sources: [E], [G], [I]; current project boundaries in [3], [10], [14].

The technical words, translated

Use the analogy to learn it, then use the technical term when it helps.

Term	Like I'm 5 / what it means here
API	A service counter: a screen sends a request and the server returns a result.
Event	A line in the notebook: an accepted thing that happened to a case.
Command	A request to add a line: it can be rejected if the action is not allowed.
State machine	Rules for which step can happen next. A card acknowledgment cannot jump directly to a pharmacy fill.
Atomic transaction	One sealed change: the relevant checks and writes succeed together or fail together.
Idempotency	Pressing the button again does not create a second copy of the same result.
Polling	Asking "anything new?" repeatedly. The live adapter uses a roughly 1.5-second interval.
Run + revision	Which notebook session, and which version of that session. They help reject old responses.
Enum	A fixed set of labeled boxes. A reason must fit one allowed category or fall back to review.
Hash + provenance	A fingerprint plus a source receipt. A hash detects changes; provenance records where the material came from.
Aggregate	A summary count. Removing names alone does not automatically establish a complete privacy guarantee.
RxFill-shaped event	A simulated message describing pharmacy fill status. The prototype does not establish a live certified network feed.

Definitions describe how terms are used in this project, based on the inspected source.

Practice until you can explain it

If you can answer these without reading the page, you understand the core story.

Five-minute rehearsal

- **Minute 1: the job.** Say the one-sentence pitch. Name the main user and the next-step problem.
- **Minute 2: the flow.** Draw six boxes from prescription to separate pharmacy confirmation.
- **Minute 3: the system.** Name the screens, server, database and polling adapter. Explain a command versus an event.
- **Minute 4: the safeguards.** Explain unknown reasons, eligibility flags, duplicate actions and source-verified label text.
- **Minute 5: the limits.** Name what is simulated, what is merely present in code, and what was physically demonstrated in the run.

Three “teach it back” questions

1. Why can't the patient tap confirm the fill?
2. Why is “AI picked the fix” an inaccurate explanation?
3. Why does a successful server notification response not prove a watch buzz?

SHORT ANSWERS

1. It is acknowledgment, not independent pharmacy evidence.
2. Rules and eligibility flags choose the administrative action.
3. Server acceptance and delivery to the physical device are different events.

Final preflight: correct deployed SHA/mode, same active run, approval ready, phone access working, label available, separate fill trigger, truthful fallback.

Use the source map on pages 31–34 to investigate a component before making a stronger claim.

Your one-page judging cheat sheet

This is the page to revisit immediately before presenting.

Remember	Say it precisely
Problem	27% in Surescripts' January 2026 study; scope: new e-prescriptions to fill-reporting pharmacies.
Main user	Coordinator owns the recurring work; the prescriber approves and handles relevant handoffs.
Core proof	Patient acknowledgment stays pending. Separate simulated pharmacy confirmation changes the fill state.
Technical contribution	Guarded commands, atomic events, run/reset isolation, polling, deterministic routing and sourced labels.
AI boundary	Gemini returns a reason or null; rules choose the fix. Null gets no reason alert and uses access support. Voice proposes a confirmed handoff.
Status update	Gemini/Tiger merged and recorded live. Frontend repairs merged; physical-device proof remains. Otezla and Humira labels verified in merged source.
Value and business	Potentially less staff effort. Pricing, partner payment and profit are hypotheses; illustrative numbers are not measured outcomes.
Do not claim	Prevented escalation, real partner fulfillment, instant fill from a tap, certified network compatibility, guaranteed savings or production compliance.

LAST LINE

"One missed fill. One accountable next step."

Cut optional scenes before cutting the disclosure or the acknowledgment/confirmation distinction. Name any switch to mock replay or recorded footage.

Sources and snapshot boundaries

This guide is based on project files and code, not a newly executed deployment or clinical evaluation.

Source links use fetched main 015c3429a12518ed67753a47db9eaae308e82b13. The deployed audit used an earlier build after PR #39 merged. This PDF reconciles evidence; it is not a new production check.

[1] Product and demo

[docs/spec-v2-coordinator.md](#)

[2] Pitch source (contains older claims)

[docs/presentation/pitch-and-qa.md](#)

[3] Current web integration handoff

[docs/handoffs/phase6-integration.md](#)

[4] Workflow authority

[lib/server/workflow.ts](#)

[5] Archived screen captures

[docs/stills](#)

[6] Patient-facing fixed copy

[mock/templates.json](#)

[7] Live polling and commands

[lib/realtime.ts](#)

[8] Verified label facade and artifacts

[lib/server/label.ts](#)

[9] Deterministic routing

[lib/server/router.ts](#)

[10] Coordinator access and message storage

[supabase/migrations](#)

[11] Notification publisher

[lib/server/ntfy.ts](#)

[12] Seeded counts and scope

[docs/handoffs/vinh-seed-week.md](#)

[13] Classifier: [lib/server/classify.ts](#) and [classifier/gemini.ts](#), merged via PR #39. Current deployed claims: [docs/claims-audit.md](#), top dated refresh; older sections are historical.

Supporting code: [app/api/*](#); [supabase/migrations/*](#); [data/labels/drug_otezla/*](#); [components/data/*](#). The guide's diagrams are explanatory illustrations. The bar chart is calculated from the repository's seeded counts; the app images are saved hosted-audit captures; see page 34.

Problem and business evidence

External facts support the problem. All FirstDose savings, prices and profit figures in this guide are scenarios.

[A] Surescripts: First-Fill Abandonment

27% of new prescriptions unfilled; footnote specifies January 2026 e-prescriptions sent to fill-reporting pharmacies. The company already sells monitoring/outreach data.

[B] AMA: 2025 Prior Authorization Physician Survey

December 2025 survey of 1,000 physicians: 40 requests per physician/week; 13 hours of physician-and-staff work; 40% report dedicated PA staff. Released in 2026.

[C] CMS: Otezla, Georgia, Part D 2024

614 prescribers and 6,661 claims. Claims include refills and are not first fills, practice counts, customers or addressable revenue.

[D] CMS: Humira(CF) Pen, Georgia, Part D 2024

717 prescribers and 15,295 claims; specific branded product. Do not add prescriber counts across drugs to infer distinct practices.

[E] DocUpdate FAQ

Staff and practice accounts are described as future roadmap work. The fill-confirmation statement is inside an answer about cancellation requests.

[F] AAMC: 2025 Key Findings

866,460 physicians in direct patient care in 2024. This source supplies physician headcount, not access-coordinator headcount or customer demand.

A PILOT TURNS A HYPOTHESIS INTO EVIDENCE

Measure net minutes and touches per case, unresolved-case age, observed fill outcomes, buyer willingness to pay and actual cost to serve. Define a comparison before claiming incremental fills or causal savings.

Additional sources and reconciliation

This master replaces the earlier guides for rehearsal. Implementation status is tied to a dated source snapshot.

[G] Impiricus product descriptions

Used for Ascend and Spark roles. These are the vendor's descriptions, not evidence of a FirstDose partnership.

[H] IQVIA historical affordability discussion

October 8, 2020 account: 44% above \$125 and >60% above \$500. Historical association; source-specific thresholds.

[I] HHS: treatment, payment and operations

Supports correcting the blanket "HIPAA forbids real data" claim. It is not an assessment of this prototype's compliance.

[14] Latest claims and deployed evidence

Use the top dated refresh; older audit sections are historical. Merged/live Gemini and Tiger evidence supersedes missing-provider claims. Latest gates and scope cuts: page 34.

What was reconciled


The story keeps Deem's coordinator-first pacing and judge participation. It now separates acknowledgment from confirmation, labels the partner concept, removes unproven clinical effects, treats payment as proposed, updates AMA survey timing, and corrects polling and Tiger status.

Inputs: Deem's supplied FirstDose-Pitch-Prep.md; the revised FirstDose PDF and business-case notes; project code/docs at main 015c3429a12518ed67753a47db9eaae308e82b13. Screens are saved hosted-audit captures preceding the frontend repairs. No fresh deployment, device, application-test or clinical evaluation was performed to produce this PDF.

All charts are either attributed external facts, explicitly synthetic fixture counts, or labeled economic assumptions. Financial arithmetic was checked; that does not validate willingness to pay or the assumed costs.

Current build and remaining proof

The pitch still fits. The implementation has moved from an integration branch to merged code and recorded deployed browser evidence.



Evidence / scope	What to say
Deployed browser audit	After #39: both workflows and independent-session persistence/reset passed; 62 route/browser/viewport checks loaded. Not 62 end-to-end cases.
Tiger summary	8 synthetic seeded fills; median 60 seconds; reason counts 2/2/2. Reset: 0 fills, null median, empty reasons. Direct hypertable query, no continuous aggregate.
Frontend repairs (#40)	Merged. Recorded 724 tests, lint/live+mock builds and 57 axe scans passed. Chromium/Firefox 44 checks each; WebKit 43 plus corrected skip-link recheck.
Voice + remaining gates	Merged pre-confirm run check; complete recording-to-approval reset safety remains unverified. Recheck latest deployed repairs; record physical flow, watch alerts, voice/audio.
Latest #43 / #44 / #47	Humira verified; owner sign-off merged. Recorded 728 tests / 54 files and local build/lint/typecheck. Five original LFS clips, not a final edit. Deployed Humira recheck remains.
Scope cuts (6.6 / 6.7)	No NPPES colleague discovery or coordinator-active/fixes-per-coordinator tiles. Existing Tiger fill/time/reason summary remains. TestFlight needs Mac acceptance.

Evidence: docs/minh-signoff-5.1.md; video/README.md; recorded hosted audit at commit 1d12279. Owner records are not fresh tests for this PDF.


Evidence is dated, not freshly rerun for this PDF. Physical phones and watches are separate from browser sessions.

Why the industry cares

The team reports that an Impiricus representative called the problem one of the industry’s biggest. Treat that as qualitative feedback on importance.

THE PROBLEM IN ONE SENTENCE

A prescription can be sent successfully while the patient still has no confirmed fill, the barrier needs attention, and the next person responsible is unclear.



Why it matters	Evidence or implication
Patient access	Surescripts reports 27% of new e-prescriptions to fill-reporting pharmacies in its January 2026 analysis were never dispensed. Not a FirstDose result.
Practice workload	AMA’s 2025 physician survey reports 13 physician/staff hours per physician per week on prior authorization. FirstDose does not automate all that work.
Operational ownership	A status, a reason, an accountable next step and a later result are different pieces of information. FirstDose connects them in one case workflow.
Partner relevance	Ascend connects HCPs to reps/resources; Spark includes prescription triggers. Our proposed fit is coordinator follow-through. Partner overlap still needs validation.

A useful distinction: importance of the problem, usefulness of the prototype, and measured impact are three different claims. The conversation supports the first; a pilot must establish the last.

Sources: [A] Surescripts; [B] AMA; [G] Impiricus. Feedback is team-reported, not a formal endorsement or an industry ranking.

Features: the coordinator's work

Explain the job first, then point to the control. Proposed value is less searching and clearer ownership; time savings are not yet measured.

1 / QUEUE, FILTERS AND WAITING-ON STATUS

Who: Access coordinator.

What: See needs-action, waiting and confirmed cases; filter the queue and see the reason, elapsed time and next owner.

Why: Turn a scattered follow-up task into a visible work list.

Boundary: The opening 3 / 2 / 8 counts are seeded fictional cases.

2 / CASE DETAIL, NEXT ACTION AND HISTORY

Who: Coordinator reviewing one case.

What: Open the case to see its reported barrier, rule-selected action, messages and event history.

Why: Keep the explanation and the next step beside the evidence.

Boundary: A rule-selected resource is an administrative action, not a treatment recommendation.

3 / PRESCRIBER LINK AND APPROVAL

Who: Coordinator requests; doctor approves in Profile.

What: The approved fictional relationship permits the case handoff into the coordinator workflow.

Why: Make delegation explicit while keeping prescribing with the doctor.

Boundary: Shared demo login and NPI format checks are not production identity verification; NPPES discovery is cut.

4 / CONTACT MARKS

Who: Coordinator making follow-up calls.

What: Mark Reached patient or Left message in the case.

Why: Leave a useful reminder of the contact attempt on that browser.

Boundary: Device-local only. Clicking a mark does not place a call, send a message or update another browser.

Say while showing it: "Here is who needs attention, why, who owns the next step, and what has already happened."

Source: current queue, case sheet, coordinator routes and local contact storage at main 015c342.

Features: the doctor's view

The doctor's surface keeps relevant status and delegation close to the prescription. It is our DocUpdate-inspired concept.

5 / NEW RX AND VERIFIED OTEZLA LABEL

Who: Doctor in the synthetic prescribing flow.

What: Record the demo order and inspect cached label sections verified against saved DailyMed source and drug identity.

Why: Put source-backed reference text beside the workflow.

Boundary: No real prescription is sent. Otezla and Humira have verified cached label artifacts. This is not an AI-generated drug label.

6 / ALERT, HANDOFF, CONCIERGE AND WATCH

Who: Doctor responding to a reported barrier.

What: Use the alert or Concierge's Help my patient start to hand off. Follow the fill line; ntfy supplies alerts.

Why: Assign the work and see the later fill signal.

Boundary: Concierge's samples/rep options are disabled. No Gemini reason means no reason alert. Server acceptance is not watch receipt.

7 / SEARCH, PATIENT DETAILS AND BEFORE-VISIT NOTE

Who: Doctor reviewing a case or preparing for a visit.

What: Find the fictional patient, history and fill status. A prepared before-visit card can flag missing confirmation.

Why: Keep unresolved access status available for review.

Boundary: The before-visit card needs a scripted/offline event; no automatic live appointment trigger. Missing confirmation is not treatment failure.

8 / OPTIONAL SPOKEN HANDOFF

Who: Doctor with a working microphone and live session.

What: Grok transcribes the supported request and proposes a case; the doctor confirms before handoff.

Why: Offer another input path for the same administrative task.

Boundary: Human/device and full recording-through-approval reset safety remain unverified. Keep the button path ready.

Say while showing it: "The doctor can see the reported problem, approve the helper, and hand over the work."

Source: doctor screens, verified label facade and ntfy path at main 015c342; device evidence remains separately dated.

Features: the patient's next step

Keep the patient experience simple. There are three different receipts; do not narrate them as one event.

9 / QR-OPENED RESOURCE PAGE

Who: Patient, or a judge playing the patient.

What: Open the case page and see the approved card/resource and current status.

Why: Make the selected next step available on the phone.

Boundary: A web resource stand-in: no SMS, Apple Wallet issuance, redeemable credential or real partner fulfillment.

10 / APPROVED ENGLISH/SPANISH MESSAGE + AUDIO

Who: Coordinator selects/approves; patient reads or listens.

What: After the scripted Otezla card-resend fix, the template and pre-generated ElevenLabs audio appear on the patient page.

Why: Offer a controlled message in two languages and an audio format.

Boundary: Not general messaging, runtime advice or Humira audio. Native pronunciation/physical playback still need review.

11 / Three receipts, three meanings

Message acknowledged

The message acknowledgment is recorded.

Resource acknowledged

The card/resource tap is recorded. Fill is still pending.

Pharmacy fill signal

A separate pharmacy event confirms the fill step.

These boxes describe different evidence, not three required sequential clicks. A message acknowledgment does not use the card; a card acknowledgment does not confirm a fill.

THE LINE TO SAY ON STAGE

“That tap records acknowledgment. Watch: the fill remains pending. Only this separate simulated pharmacy confirmation changes the fill state. Even a fill does not prove a dose was taken.”

Source: patient page, persisted message APIs, templates and pharmacy simulator at main 015c342.

Features: intelligence and evidence

Name each technology only after explaining its job and why that job matters.

12 / GEMINI REASON CLASSIFICATION

Who: The workflow, on a supported pharmacy/hub note.

What: Map the fictional source note into one allowed reason code or null.

Why: Make reported barriers usable by a consistent workflow.

Boundary: No invented reason on failure; a four-second timeout becomes null. AI does not select the fix.

13 / DETERMINISTIC RESOURCE ROUTING

Who: Server, followed by coordinator review/action.

What: Use the reason and explicit eligibility flags to choose a permitted resource path.

Why: Keep administrative choices repeatable and constrained.

Boundary: Unknown/ineligible cases use access support; simulated programs do not prove real eligibility or fulfillment.

14 / TIGER FILL, TIME AND REASON SUMMARY

Who: The proposed access-program/buyer view.

What: Read run-scoped confirmed-fill counts, median time to first fill and reason counts.

Why: Let reviewers inspect observed workflow results, not just clicks.

Boundary: Synthetic results are not impact or ROI. Direct hypertable query; no continuous aggregate or coordinator-activity tiles.

15 / PUBLIC MARKET CONTEXT

Who: Team or prospective buyer exploring rollout.

What: Show sourced Georgia prescribing/formulary reference data separately from demo case metrics.

Why: Ground a market discussion in named datasets rather than invented demand.

Boundary: Claims, prescribers, practices and customers are different units. These data do not establish our addressable market.

16 / Shared state and safeguards. Supabase commits guarded events; 1.5-second polling refreshes screens. Run/revision checks and duplicate handling protect consistency. Tiger receives allowlisted HMAC metrics; the demo login is not a production role/privacy system.

ELI5: Gemini reads the label on a problem; the rulebook chooses an allowed task; the notebook remembers it; the scoreboard counts later evidence.

Sources: classifier, router, analytics, public reference data and workflow modules at main 015c342; [C/D] CMS for market context.

Explain every feature without rushing

Completeness means every feature has a clear job. Stage time belongs to the core workflow; optional paths and engineering details belong in questions.

USE THIS SENTENCE FOR EACH FEATURE

“For [person], this [feature] shows or does [specific action], so they can [next step]. Here is what the demo proves.” Then point to the screen and perform the action.

Feature group	Short demonstration line
Coordinator queue	“This is the work list: who needs attention, why, and who acts next.”
Doctor + source label	“The order starts the workflow; the Otezla reference text has a saved source.”
Reason + rule + handoff	“Gemini reads the reported reason. Rules constrain the resource. The doctor delegates; the coordinator acts.”
Patient message/resource	“The patient gets the approved resource and, for Otezla, a message they can read or hear.”
Acknowledgment vs fill	“The tap stays pending. This separate simulated pharmacy event confirms the fill.”
Tiger summary	“Now count the recorded fills and time in this synthetic run; these are not pilot outcomes.”

17–20 / Explain the demo controls too

Setup/QR: prepares the right screens. **Seed/reset:** loads a repeatable fictional week or starts a clean run; reset clears links/messages.

Simulator + raw RxFill-shaped view: supplies labeled pharmacy evidence, not a live pharmacy feed. **Board:** optional timeline display for the audience; not another user product.

Keep optional voice, before-visit detail and public-market tiles ready for questions. Do not promise every feature as a separate live scene in a four-minute pitch.

Feature scope: main 015c342. NPES discovery and coordinator-specific activity tiles are cut. Partner services remain simulated.

Khadim · your speaking script

Open, introduce the user and product, hand over to Vinh, then close. Lead business questions.

0:00–0:15 / The problem

“DocUpdate's headline says: ‘The prescription was sent. The patient still never started it.’ Surescripts found 27 percent were never dispensed in its analysis. Sending a prescription does not finish the job.”

Action: Face the judges; keep the watch ready.

0:15–0:30 / The person

“Impiricus reports over a million opted-in healthcare providers. We focus on another person in the practice: the access coordinator. Their job is helping patients get the medicine the doctor prescribed.”

Action: Point to the coordinator desktop without starting the demo.

0:30–0:50 / What it is

“FirstDose is our proposed Ascend skill for DocUpdate: a staff workflow after a reported fill barrier. Spark engages on signals; FirstDose turns a barrier into an owned task. We would test Market Access payment per qualifying confirmed first fill. Vinh, show them.”

Action: Hand the floor to Vinh.

2:45–3:05 / Three proposed upgrades

“We build on the Impiricus ecosystem. DocUpdate gets a proposed fill signal and staff seat. Ascend gets a barrier-to-task workflow. Wallet, Concierge, QPharma and Medvantx are the resource destinations. Our next proof is a practice pilot measuring staff time and independently confirmed fills.”

Action: Take the floor from Vinh. Point to the ecosystem diagram if time allows.

3:05–3:15 / Close

“Impiricus reaches the doctor who writes the prescription. FirstDose reaches the person who gets it filled. One missed fill. One accountable next step.”

Action: Stop. Khadim takes business, Vinh build, Minh AI and label questions.

WALK-UP / ABOUT 10 SECONDS

“Sending a prescription does not finish the job. FirstDose connects a reported barrier to the coordinator who can act. Put this on.”

Questions: Khadim — buyer, market, partner fit. Vinh — build, data flow, polling, watch. Minh — AI and labels; Vinh supports. Keep optional features out if they risk the core proof.

Same spoken lines as master pages 4–5. Times are rehearsal targets; read actions silently.

Vinh · your speaking script

Drive the coordinator, doctor and patient demo. Lead data-flow, build and watch questions.

0:50–1:00 / Her morning

“This queue has three needing a fix, two waiting, and eight filled. These are fictional cases. Supabase stores the shared state; polling keeps the screens synchronized.”

Action: Show the seeded queue before adding interactive cases.

1:00–1:10 / The prescription

“This is our DocUpdate-style prescribing demo. The Otezla label is verified against saved DailyMed source. New FirstDose tags mark our proposed additions.”

Action: Sign the synthetic order and show the label.

1:10–1:25 / It breaks

“The simulated pharmacy reports a barrier. Gemini reads the note into a reason: declined at price. Read the alert to me.”

Action: Trigger barrier/reason; invite the judge to read the received alert. If null, use Patient Details/access support.

1:25–1:35 / One tap

“Send to my coordinator. One handoff tap after setup. The doctor first approves the coordinator relationship. Optional voice proposes the same handoff for confirmation.”

Action: Use the approved link, or show the approval step. Use voice only if rehearsed.

1:35–1:45 / The fix

“A rule picks the fix, not AI. Government coverage never gets a manufacturer copay card in our router. The coordinator reviews and sends through our Wallet stand-in.”

Action: Send the permitted resource. Prepare the approved message if using audio.

1:45–2:15 / The judge plays the patient

“Now you're Maria. Tap Use at pharmacy. That tap is acknowledgment; we still wait for the pharmacy. Now our operator sends a separate simulated fill confirmation. The board updates. The doctor gets two alerts: the barrier and the confirmed fill. Minh built the AI part.”

Action: Hand over QR/prepared phone. Minh fires confirmation only on cue after acknowledgment. Report any missing wrist alert honestly.

2:35–2:45 / The proof

“The rule routes to access support. Tiger receives selected metrics: confirmed fills and time to first fill. This view shows aggregates without names.”

Action: Show summary with its actual source. If fallback, say practice event counts.

Questions: Khadim — buyer, market, partner fit. Vinh — build, data flow, polling, watch. Minh — AI and labels; Vinh supports. Keep optional features out if they risk the core proof.

Same spoken lines as master pages 4–5. Times are rehearsal targets; read actions silently.

Minh · your speaking script

Speak for about 20 seconds in short sentences. Lead AI and label questions; Vinh supports.

2:15–2:35 / James — short sentences

“James Carter: Humira first fill confirmation is still pending.

Gemini reads the hub note.

It gives one reason: unable to reach.

If it is not sure, it leaves the reason unknown.

Gemini does not pick the fix.

A rule does.”

Action: After the opening queue, prepare James: sign ev_14, then fire ev_16–18 in order. At 2:15 Vinh shows the result; if null, explain that fallback instead.

YOUR BEST AI ANSWER

“AI reads the note. A rule picks the fix. A human taps send.”

Know these three follow-ups

- **Why not prior auth?** The earlier pharmacy reject is 75. Gemini classifies the later hub contact note, whose expected reason is unable to reach.
- **What does unknown mean?** The backend returns null for uncertain, invalid, rejected or timed-out classification. It does not invent a reason.
- **What changed?** Otezla and Humira now have verified cached labels. Humira does not gain audio or a live visit trigger from that change.

Questions: Khadim — buyer, market, partner fit. Vinh — build, data flow, polling, watch. Minh — AI and labels; Vinh supports. Keep optional features out if they risk the core proof.

Same spoken lines as master pages 4–5. Times are rehearsal targets; read actions silently.

What changed in this revision

Your new speaking order is preserved. These corrections make the same story hold up during questioning.

Supplied line / topic	Use this version
Speaker ownership	Khadim opens/closes and takes business. Vinh drives the demo and build questions. Minh has a short AI scene and takes AI/label questions.
Supabase Realtime	Supabase stores shared state; authenticated HTTP polling about every 1.5 seconds keeps the views current.
James = prior auth	Earlier pharmacy reject 75 is prior auth. Gemini reads the later hub note about failed contact, expected UNABLE_TO_REACH. Rule: ACCESS_SUPPORT.
Every event / counts only	Only allowlisted, reduced metric events reach Tiger. The view shows aggregates. HMAC identity is pseudonymous; shared demo login is not production buyer isolation.
First staff seat / Wallet	Proposed staff workflow and partner integration. Wallet/QPharma/Medvantx fulfillment remain stand-ins; no claim that Impiricus has never served staff.
New code and media	Both labels verified; Humira includes boxed warning and stronger identity checks. Minh sign-off and five original LFS source clips merged, not a final edited film.
Fill / before visit	Acknowledgment is not fill. Simulated pharmacy confirmation is separate. Before-visit card needs a prepared scripted event; there is no live visit trigger.

Pricing per qualifying confirmed fill is a hypothesis. “Industry’s biggest problem” is team-reported qualitative feedback, not a ranked study or company endorsement.

Code snapshot: 015c342. No fresh application, deployed Humira or physical-device test is claimed by this document.

Sources behind the new script

Checked September 26, 2026. Preserve the population, date and uncertainty when quoting a number.

J / DocUpdate headline

The supplied headline is real; the homepage dates the article July 9, 2026, by Sana Khateeb, PharmD. It establishes published problem recognition, not a FirstDose endorsement.

K / Impiricus reach

The company's own job posting says more than one million opted-in HCPs. HCPs is broader than doctors. This is company-reported network reach, not FirstDose customers or active DocUpdate users.

L / Workforce context

Current page reports 833,900 medical-assistant jobs in 2025. The proposed 467,000 office estimate needs its 56% share independently sourced. Neither number measures access coordinators or customers. Removed from the timed pitch.

M / Latest label and owner evidence

Humira PR #43 is now merged. Recorded branch-local label endpoint/render checks and 728 tests / 54 files; deployed Humira recheck remains distinct. Gemini pin and Tiger parity evidence are recorded here.

N / Original demo clips

Five source recordings: Khadim beats 1–3, Minh beat 1, Vinh beat 1. Presence is not verification of spoken claims, a final edited submission or physical-device acceptance.

WHY THE PROBLEM MATTERS

The evidence supports a substantial first-fill gap and recurring administrative work. FirstDose proposes a coordinator workflow around it. Savings, incremental fills, partner distribution and willingness to pay still need a pilot.